

Intro To Machine Learning – Assignment-3

UB PERSON NUMBER: 50425014

UB IT NAME: PARAVAMU

TASK: The goal of the assignment is to learn the trends in stock price and perform a series of trades over a period and end with a profit. In each trade you can either buy/sell/hold. You will start with an investment capital of \$100,000 and your performance is measured as a percentage of the return on investment.

You will use the Q-Learning algorithm for reinforcement learning to train an agent to learn the trends in stock price and perform a series of trades. You will implement Q-learning algorithm from scratch. The purpose of this assignment is to understand the benefits of using reinforcement learning to solve the real-world problem of stock trading.

STEPS INVOLVED:

1.) IMPORTING NECESSARY LIBRARIES

- ⇒ Importing numpy and pandas for data processing
- ⇒ Importing Gym for developing our environment for the reinforcement problem.
- ⇒ Importing matplotlib to show images
- ⇒ Importing random for randomizing actions

```
[6] 1 # Imports
    2 import gym
    3 from gym import spaces
    4 import math
    5 import matplotlib.pyplot as plt
    6 import numpy as np
    7 import pandas as pd
    8 import random
```

2.) IMPORTING THE DATASET

```
1 stock_trading_environment = StockTradingEnvironment('./NVDA.csv', number_of_days_to_consider=10)
2
```

- ⇒ The dataset given to our environment is the historical stock price for NVIDIA for the last 5 years
- ⇒ The features include information such as: Date, Open, High, Low, Close, Adj Close, Volume.
- ⇒ We split the data into train and test data
- ⇒ We use the train set for agent training and we use the test set for agent evaluation

3.) STOCK TRADING ENVIRONMENT

⇒ Our stock trading environment has three main functions.

⇒ They are:

- Reset ()
- Step ()
- Render ()

⇒ Reset ()

- This method resets the environment and returns the initial observation
- There are 4 types of states or Observations ranging from 0 to 3
- Based on the observation vector it returns the observation

```
110
111     # Observation vector.
112     observation = [price_increase, price_decrease, 0, 1]
113     if np.array_equal(observation, [1, 0, 0, 1]):
114         observation = 0
115     if np.array_equal(observation, [1, 0, 1, 0]):
116         observation = 1
117     if np.array_equal(observation, [0, 1, 0, 1]):
118         observation = 2
119     if np.array_equal(observation, [0, 1, 1, 0]):
120         observation = 3
121
122     return observation
```

⇒ Step ()

- The step method takes an action as the input.
- There are three actions for us to consider
- Buy, Sell and Hold
- Based on the action given it returns the next Observation, Reward, If the episode is done or not and some info

```
336     return observation, reward, done, info
337
```

⇒ Render ()

- The render method is a function which uses matplotlib to plot the total account value over time.

```
342
343
344     plt.figure(figsize=(15, 10))
345     plt.plot(self.total_account_value_list, color='lightseagreen', linewidth=7)
346     plt.xlabel('Days', fontsize=32)
347     plt.ylabel('Total Account Value', fontsize=32)
348     plt.title('Total Account Value over Time', fontsize=38)
349     plt.grid()
350     plt.show()
```

4.) QLEARNING IMPLEMENTATION

- ⇒ QLearning learns the underlying value of the possible actions in a particular observation or state.
- ⇒ Its model free unlike other reinforcement algorithms.
- ⇒ It also takes a discount factor so it can determine the importance of the future reward.

5.) INITIALIZING THE PARAMETERS

- ⇒ We'll pass the stock trading environment as a parameter, so we can use all the methods provided by the environment

```
2 Agent = QLearning(stock_trading_environment)
```

- ⇒ Number of episodes denotes how many epochs the agent should train for.
- ⇒ We'll need the Qtable initialized with zeros so we can use it to fill the rewards and predict our actions.
- ⇒ We'll need all the parameters for making our epsilon decay algorithm.
- ⇒ Also, we'll need two lists to track the reward dynamics and epsilon decay.
- ⇒ We'll need to also initialize the possible actions, so we can later randomly pick an action to feed to our step function.

```
10 self.environment = environment
11 #no of epochs it should run for
12 self.numOfEpisodes = 500
13 #stores the epsilon decay over time
14 self.exploreProbabilityArray = []
15 #stores the reward over time
16 self.rewardsArray = []
17 #initializing the Qtable with zeros states x actions
18 self.Qtable = np.zeros([4,3])
19 self.epsilon = 0.9
20 self.epsilonMin = 0.00001 # minimum exploration probability
21 self.epsilonDecay = 0.000005 # exponential decay rate for exploration prob
22 self.initialState = self.environment.reset()
23 self.action = [0,1,2]
24 self.decayStep = 0
25 self.exploreProbability = 1;
```

6.) TRAINING OUR AGENT

- ⇒ The training runs the number of episodes given above.

```
42
43     for episode in range(self.numOfEpisodes):
44
```

- ⇒ We get the initial state with the help of environment reset function.

```
45     #gets the initial state from the environment
46     currentState = self.environment.reset()
```

- ⇒ To denote an episode completion, the step function returns done.
- ⇒ So, till we get a done, we will run one episode.
- ⇒ About the Qtable, the Qtable can be formed with the help of Bellman Equation

```
57 #we use bellmans equation to form the qtable. with this we can get the actions
58 self.Qtable[currentState,possibleAction] = nextReward + (self.epsilon*np.max(self.Qtable[nextState,: ]))
```

- ⇒ The Qtable will help us predict the actions.
- ⇒ But we don't know what action to start from (The action to be passed to the step function to get our next observation and its corresponding reward).
- ⇒ So we randomize the actions initially with the help of random.choice.
- ⇒ With the help of this we complete every episode and complete our Qtable.
- ⇒ But this is not enough. We are exploring 100% to complete the Qtable and this is not very efficient and not the expected result.
- ⇒ To make this efficient we use the Epsilon decay algorithm.
- ⇒ We make the nextState as the current State and run the algorithm till we get a done from the step function and complete all our episodes.
- ⇒ We append the cumulative reward for each episode also append the epsilon decay change over time.

```
72 self.rewardsArray.append(cumulativeReward)
73 self.exploreProbabilityArray.append(self.exploreProbability)
```

- ⇒ The training iterations for few episodes are shown below:

```
[43] 1 Agent.train()
      2
      Episode Number: 0
      Epsilon decay: 0.8959752709752474
      Rewards: -2747.9922532831242
      Episode Number: 1
      Epsilon decay: 0.8919685404212803
      Rewards: -2455.783088165446
      Episode Number: 2
      Epsilon decay: 0.8879797278494629
      Rewards: -2506.967828303273
      Episode Number: 3
      Epsilon decay: 0.8840087531311016
      Rewards: -2660.7703487069853
      Episode Number: 4
      Epsilon decay: 0.8800555364958373
      Rewards: -2198.2140107772584
      Episode Number: 5
      Epsilon decay: 0.8761199985300405
      Rewards: -2135.4274699228527
```

7.) EPSILON DECAY ALGORITHM

- ⇒ The epsilon decay algorithm controls two things:
 - How much should the algorithm explore.
 - How much should the algorithm exploit.
- ⇒ Exploitation is possible only if the exploration is done and tracked properly.
- ⇒ The act method contains the epsilon decay algorithm.
- ⇒ It has two choices based on the constraint.

- ⇒ It checks if the probability which slowly decreases over time is greater than `random.rand` value (between 0 to 1).
- ⇒ We usually start with 1: which means that we ask the algorithm to explore 100% times.
- ⇒ But this slowly decreases over time and the probability of exploration decreases and we pick the argmax reward from the qtable [nextState row] to get the best possible action.

```

27 def act(self, decayStep, currentState, episodeNumber):
28     self.exploreProbability = self.epsilonMin + (self.epsilon - self.epsilonMin) * np.exp(-self.epsilonDecay * decayStep)
29     # print("exploreProbability:", self.exploreProbability)
30     randInt = np.random.rand();
31     if self.exploreProbability > randInt:
32         # print("rand called")
33         return random.choice(self.action), self.exploreProbability
34     else:
35         # print("arg called")
36         return np.argmax(self.Qtable[currentState,:]), self.exploreProbability
37

```

8.) EVALUATING OUR ALGORITHM


- ⇒ Now that the Qtable is completed, we can pick the best actions, observations and get the maximum total account value over time.

```

78 def evaluate(self):
79     """This method evaluate the trained agent's performance."""
80     """TO DO: Evaluate the trained agent's performance by selecting only the greedy/best action in each state."""
81     self.environment.train = False
82     #running it over test data
83     self.environment.reset()
84     for episode in range(1):
85         currentState = self.environment.reset()
86         done = False
87         while not done:
88             #picking actions based on qtable
89             possibleAction = np.argmax(self.Qtable[currentState,:])
90             nextState, nextReward, done, info = self.environment.step(possibleAction)
91             currentState = nextState
92             #to display the total account value over time
93             self.environment.render()

```

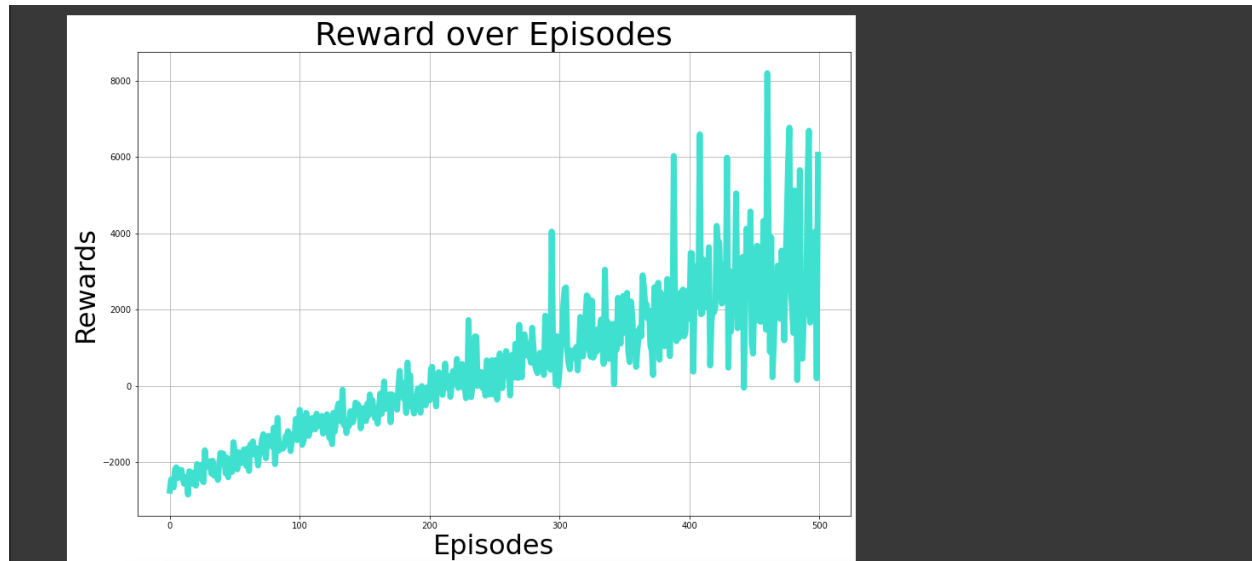
- ⇒ We call the render method to view our results.



9.) PLOTTING OUR RESULTS

⇒ With the help of the plot function, we can visualize the reward dynamics over time

```
107 plt.figure(figsize=(15, 10))
108 plt.plot(self.rewardsArray, color='turquoise', linewidth=7)
109 plt.xlabel('Episodes', fontsize=32)
110 plt.ylabel('Rewards', fontsize=32)
111 plt.title('Reward over Episodes', fontsize=38)
112 plt.grid()
113 plt.show()
```



⇒ We also plot the epsilon decay over the episodes

